Beginner-friendly ML school

Viviana Acquaviva
(CUNY)
vacquaviva@citytech.cuny.edu

What is machine learning?

THE ART OF TEACHING A MACHINE TO MAKE DECISIONS

(e.g. recognize objects, similarities and differences, patterns, signal vs noise)

Our brain machine learns

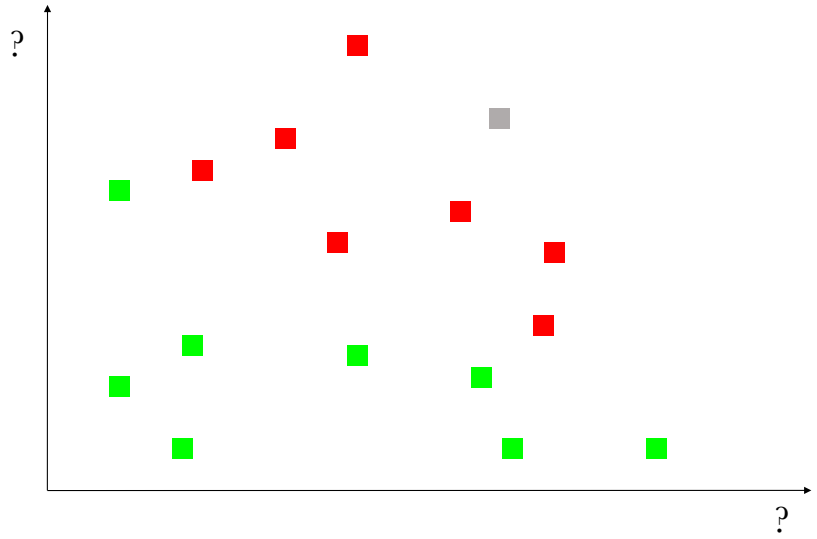

(of course)
ML is not the only way

I CAN ALSO WRITE A FORMULA
(MAKE A MODEL)
TO PREDICT COLOR BASED ON COORDINATES

MACHINE LEARNING ALGORITHMS
PROVIDE AN "IMPLICIT MODEL" IN THE
FORM OF A PATH TO A DECISION
AND PERHAPS RESEMBLE MORE THE WAY
WE (HUMANS) SOLVE PROBLEMS

MACHINE LEARNING JARGON

Features are observable quantities known for all objects

Label is the target property that we want to predict

SUPERVISED ML ASSUMES THAT WE HAVE A SET OF OBJECTS WITH KNOWN LABELS, called the LEARNING SET

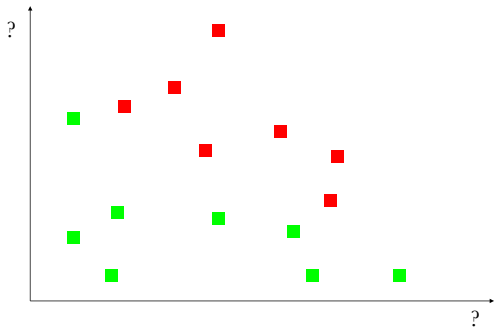

Performance is limited by size and quality of the learning set.

0 , 0
2 , 2

3 , 3 ?

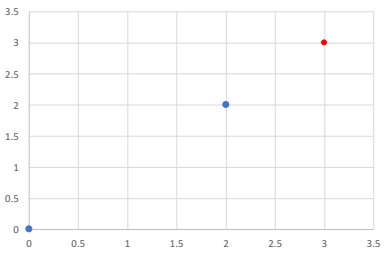

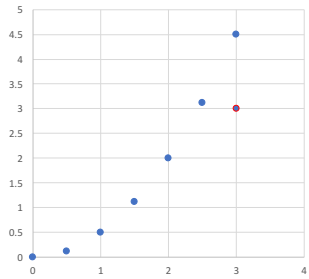

Data representation (sometimes called feature engineering) and determining whether you have enough data to create a good model are crucial.

UNSUPERVISED MACHINE LEARNING

No labeled examples

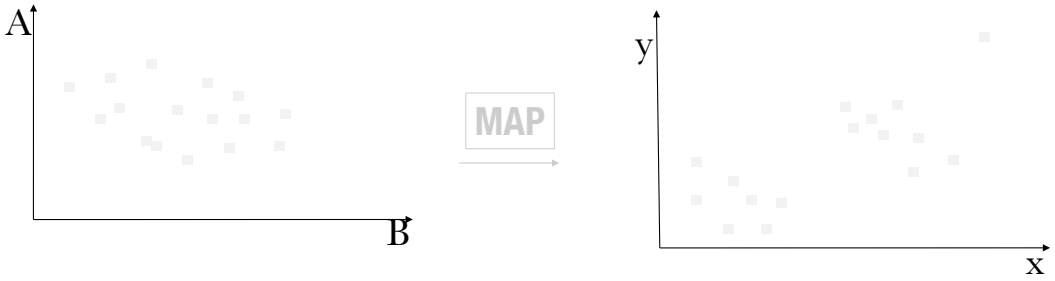

Useful to group together similar objects, find outliers, or find more efficient representations of data

Can be combined with human input or limited labels to understand the groups

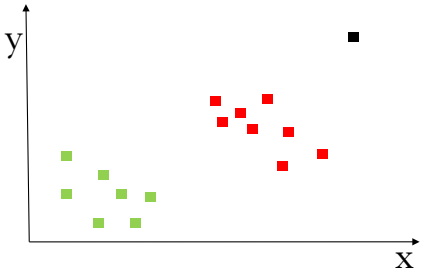

Regression vs. classification

Usually we talk about classification when the target is a discrete variable (or class). For example in this image recognition problem:

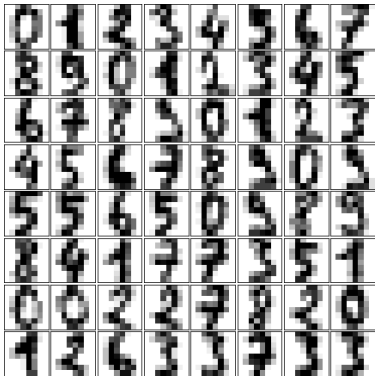

There are a finite (10) numbers of possible outcomes.

Regression vs. classification

- Vice versa, if we are trying to predict, say, the probability that it will rain in half a hour based on the current weather conditions, the outcome (target) is a continuous variable that can have all values between 0 and 1.

What if I was trying to decide whether or not I should bring an umbrella?

ALGORITHMS ABOUND

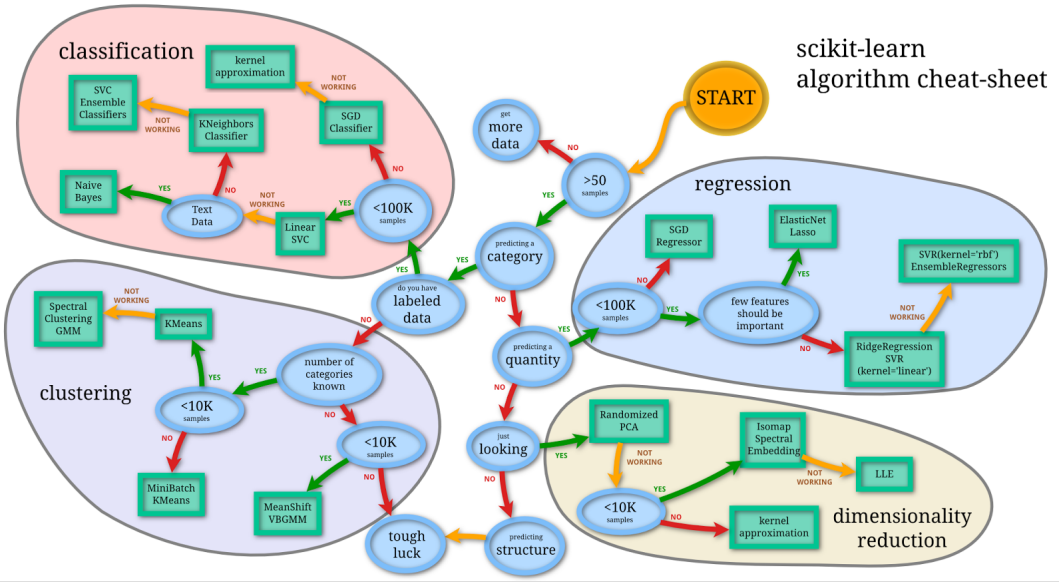

What can we do with them?

1. Save time.

Galaxy morphology

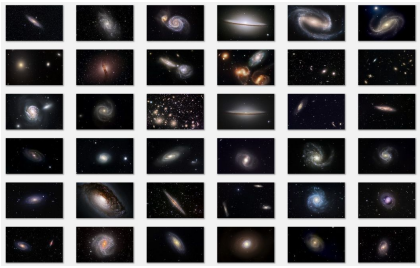

Trained humans are the best classifiers.
But what to do when
you have millions of objects?

Citizen science

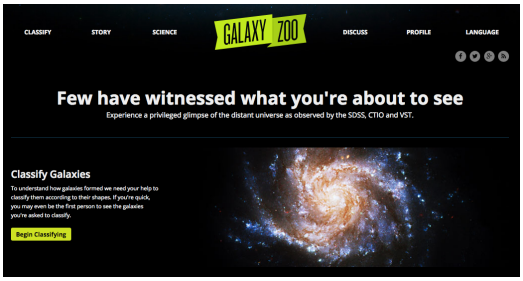

Automated Classification via Machine Learning

(supervised/unsupervised approach, see e.g. Hocking et al 2017)

2. Provide an alternative to simplistic models.

≠

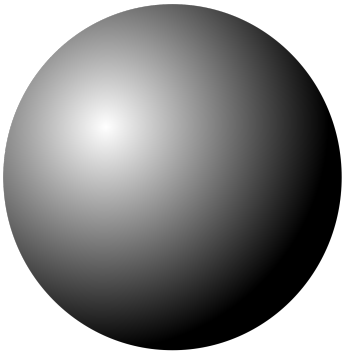

2. Provide an alternative to simplistic models.

Example: Galaxy Photometric Redshifts or SFH.

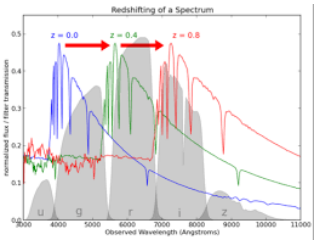

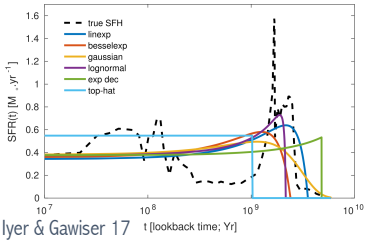

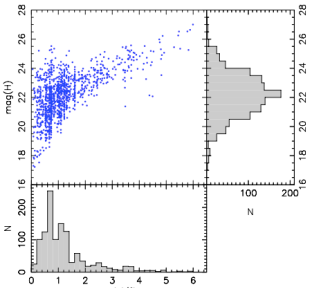

Dahlen et al 2013
Sadeh et al 2017 (AnnZ2)

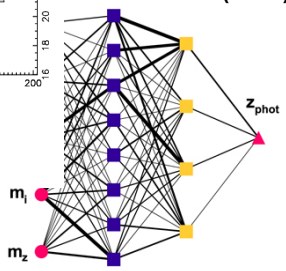

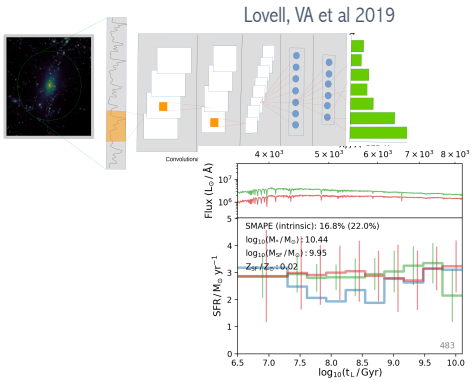

3. Make problems more tractable, e.g.:

Via dimensionality reduction

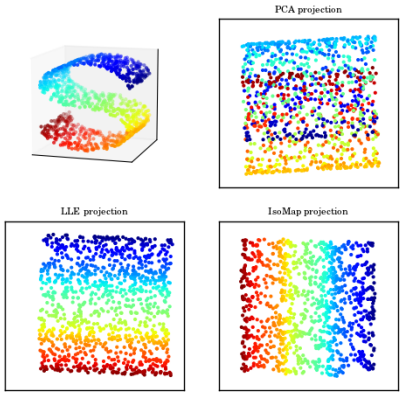

3D -> 2D
With various degrees of information loss

By making mix-and-match easier

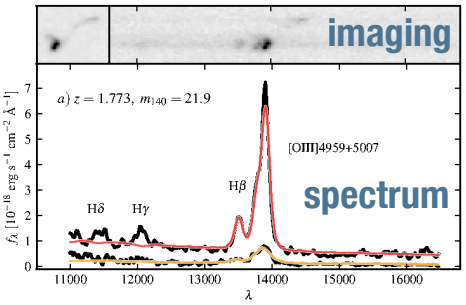

Parameter estimation:
P(model | combined data) = mess ☹

In machine learning:
Higher tolerance to mixed features ☺

4. Allow serendipitous and data-driven discoveries

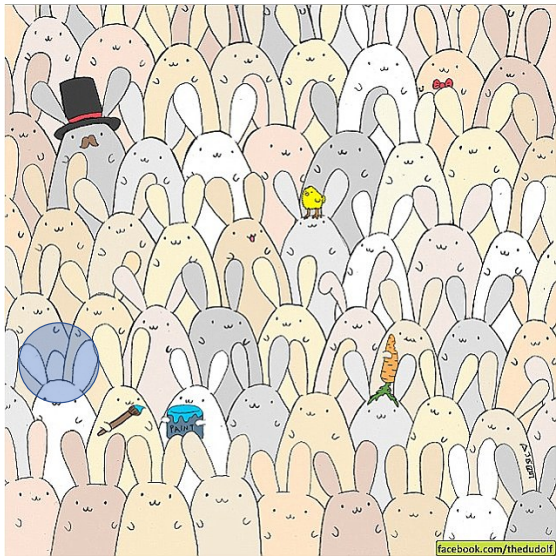

Outliers might be very interesting objects. In model fitting they are often discarded. Data mining techniques can tell you about new categories of objects.

Machine Learning vs Model Fitting

ML	MF
• Data-driven (only as good as the data)	• Intuition or model-driven (only as good as the scientist :))
• Usually generalizes poorly (model derived using some data can't be applied blindly to different data)	• Generalizes well if model (physics) is well understood
• Answers questions; interpretation is possible but might be non-trivial	• Builds subject matter knowledge; easier to interpret
• Fast(er)	• Might be computationally intensive
• More robust/accommodating of mixed and missing data	• Dealing with heterogenous data often a pain in the neck
• Allows serendipitous discoveries	• Leads to loss of information if models are too simplistic

Synergy is often the best strategy

Questions?

Let's play a fun game ☺

Let's answer these questions

Is this supervised or unsupervised ML?

Is this a classification or regression problem?

What could be useful features (data) to collect?

Let us focus on supervised learning for now.

Important:
you shouldn't use all of your learning set to build your model.

It is customary to split the learning set into a training set and test set

By building a model on the training set and applying it to the test set, you "mimic" what happens when your model sees new data for which the labels are not known.
Otherwise you would be too optimistic!

Note:
You still are.

For example…

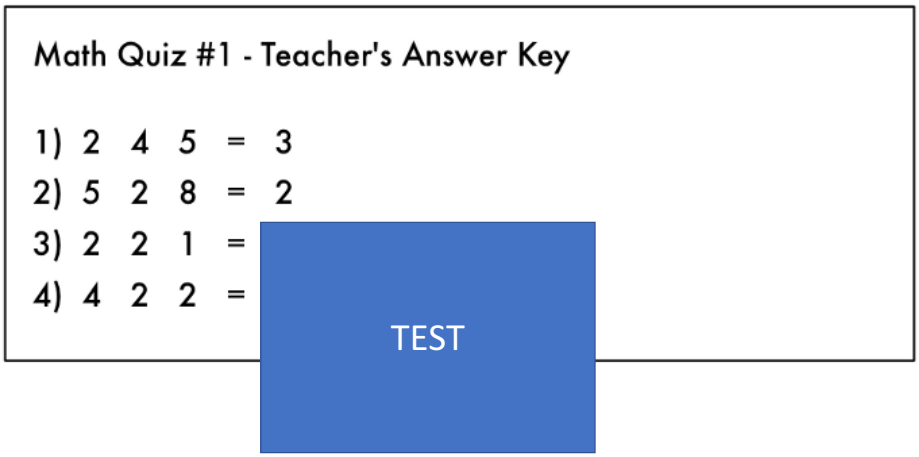

Diagnosing and Improving Machine Learning algorithm performance: cross validation, performance metrics, diagnostics

The goal of the training set and test set split is to be able to evaluate performance on unseen examples.

The test set "mimics" new data.

However, it might be better to pick more than one test split.
Why would this be a good idea?

1. We use all the training data for training!

2. We avoid the risk of under/overestimating performance because of a "weird" pick of train/test split.

k-FOLD CROSS VALIDATION

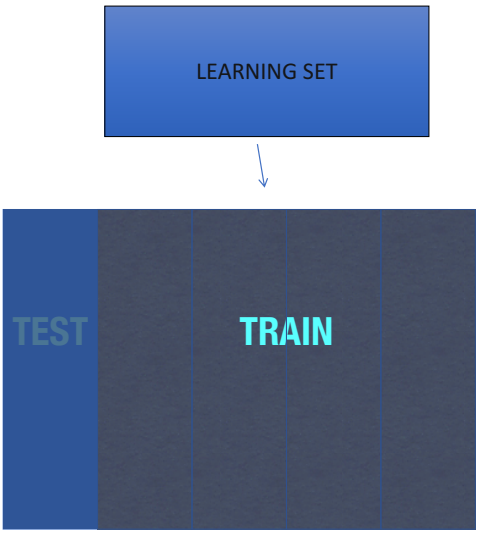

5 folds (k = 5)

k-FOLD CROSS VALIDATION

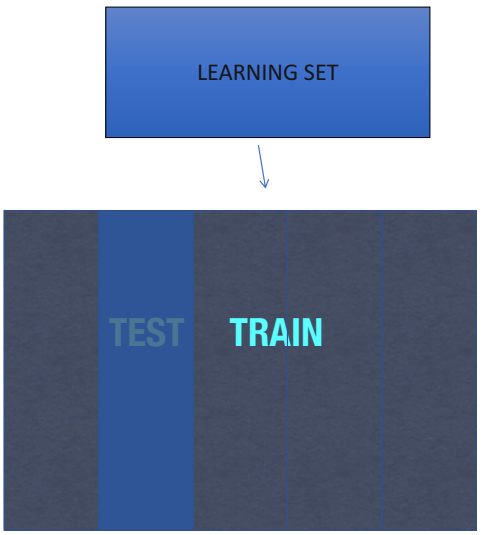

5 folds (k = 5)

k-FOLD CROSS VALIDATION

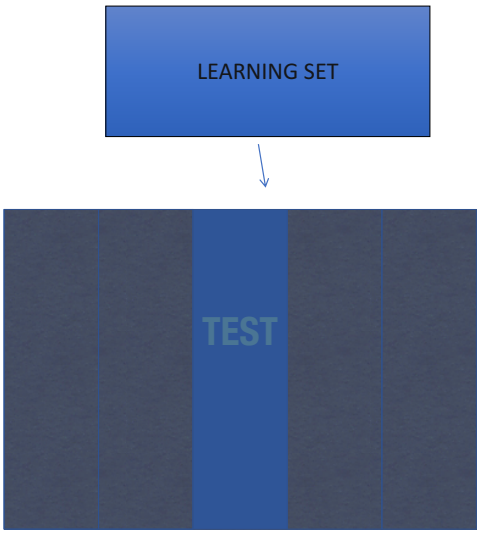

5 folds (k = 5)

k-FOLD CROSS VALIDATION

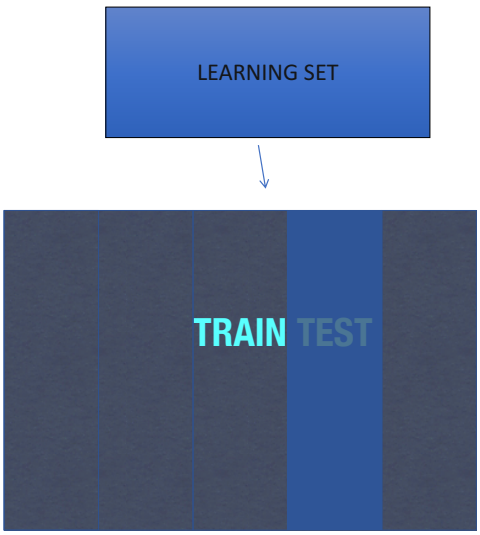

5 folds (k = 5)

k-FOLD CROSS VALIDATION

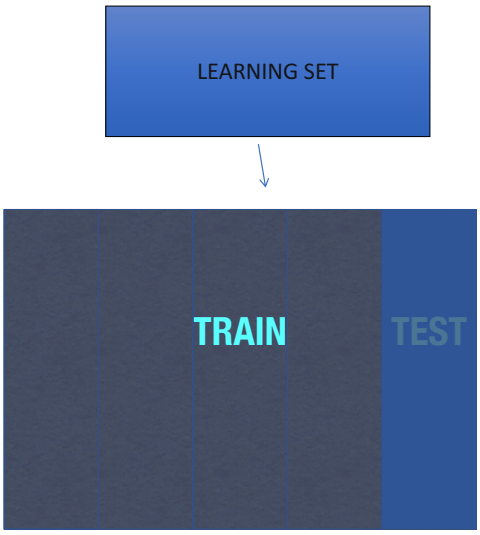

5 folds (k = 5)

k-FOLD CROSS VALIDATION

☺ We use all learning data

☺ The standard deviation of scores vector gives us an idea of average performance + uncertainty

☹ It takes 5x more time to run

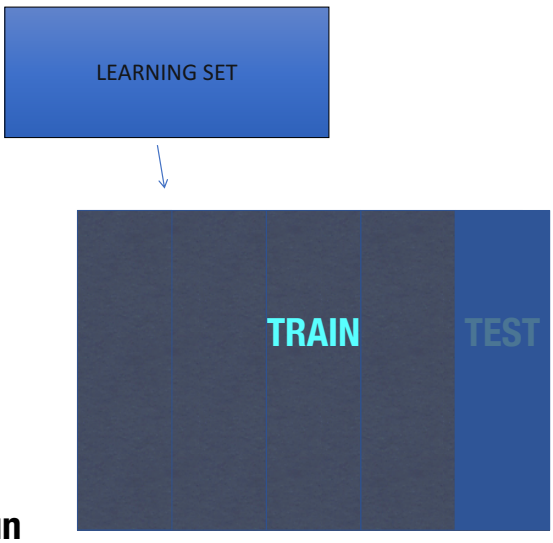

5 folds (k = 5)

Diagnosing a ML algorithm

BIAS

Algorithm
can't capture
complexity
of rule
connecting
input and output

VARIANCE

Algorithm
is excessively
tailored
to training set
and generalizes
poorly

UNDERFITTING

OVERFITTING

Bias/variance

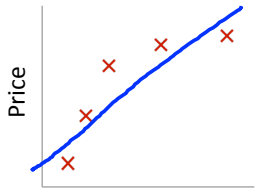

$\theta_0 + \theta_1 x$

High bias
(underfit)

$d=1$

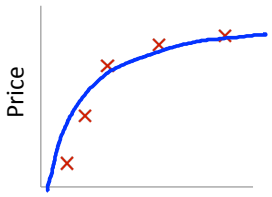

$\theta_0 + \theta_1 x + \theta_2 x^2$

"Just right"

$d=2$

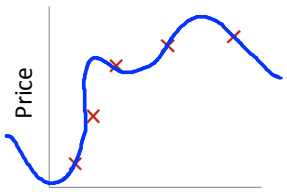

$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$

High variance
(overfit)

$d=4$

Andrew Ng

slide from Andrew Ng's Coursera ML class

How can we diagnose high variance vs high bias?

figure from Jake VanderPlas' book

High bias: test and train error are similar but high

High variance: there is a gap between test and train error because algorithm does not generalize well

Improving high bias

1. Try using different features.
2. Try engineering new features.
3. Try a more complex algorithm.

Improving high variance

1. Try reducing the number of features.
2. Try a less complex algorithm/change parameters.

Also: Check if you need more training data

Useful diagnostics: learning curves

plot performance of algorithm for train and test set as a function of size of training set

Note: which algorithm is the best?

- High bias low variance
- High variance low bias
- Lowest gap between train/test
- Highest test scores

Note: which algorithm is the best?

- High bias low variance
- High variance low bias
- Lowest gap between train/test
- Highest test scores

DECISION TREES

- Work by splitting data on different values of features
- Simplest trees are binary trees
- If categorical features, the split would be on yes/no
- If numerical, the split would be on a certain value (e.g. x > 100 or x < 100)

Example: Look at this 2-feature data set. How should we split?

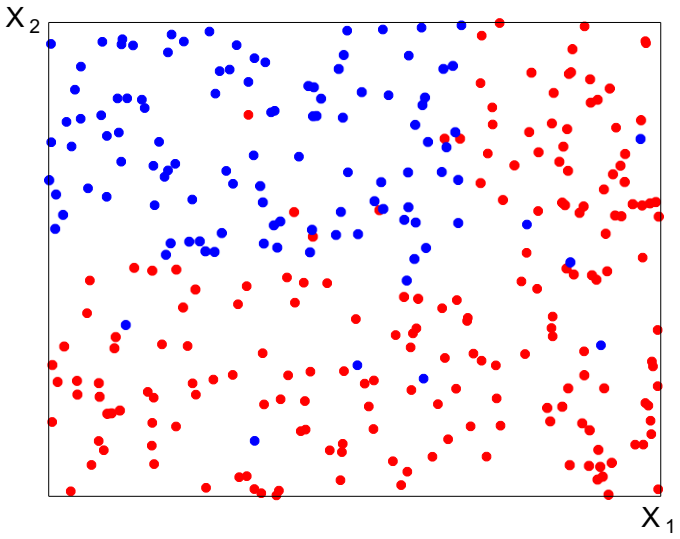

Figure credit:
Gilles Louppes

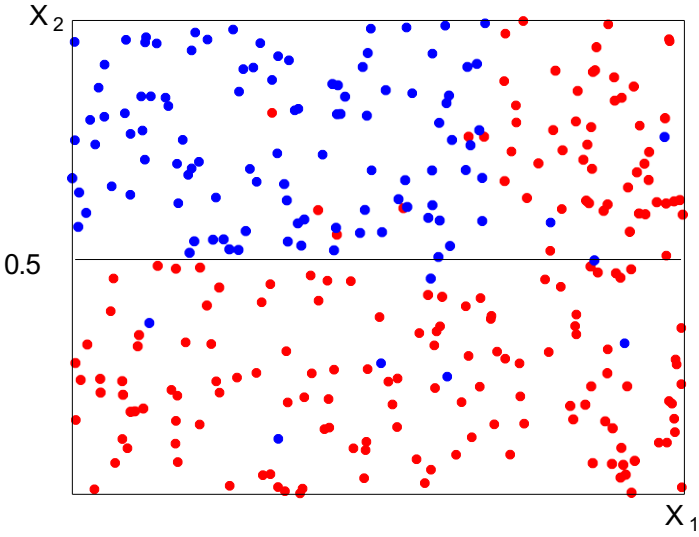

Figure credit:
Gilles Louppes

Should we stop?

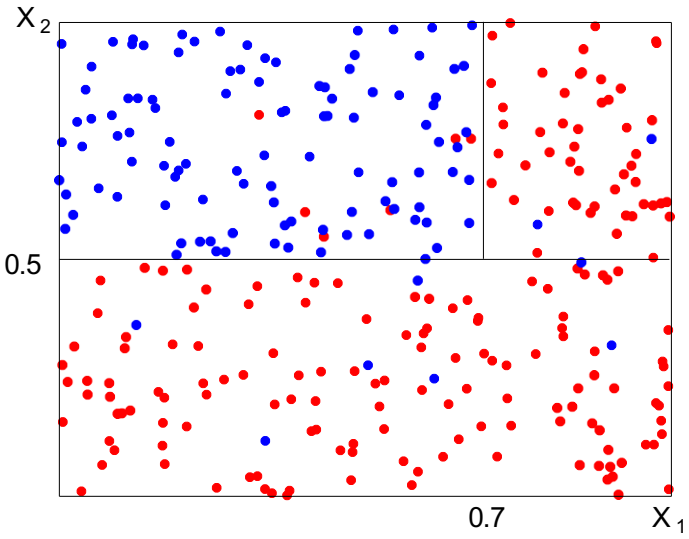

Figure credit:
Gilles Louppes

Decision trees: defined by splits and leaves

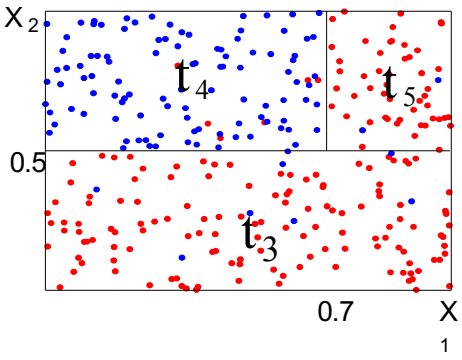

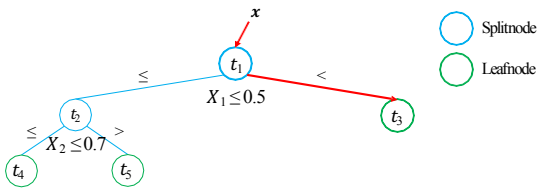

Figures credit:
Gilles Louppes

How many splits in this tree?
How many leaves?
How do we decide whether we should keep splitting?

Pseudo code for decision trees

```
function BuildDecisionTree(L)
   Create node t from the learning sample Lt = L;
   calculate (im)purity
   if the stopping criterion is met for t then
     y^ = some constant value/class  (MAKE PREDICTION)
   else
      Find the split on Lt that maximizes impurity
      decrease
       s* = arg max Δi (s, t)
                    s∈Q
      Partition Lt into LtL ∪ LtR according to s*
            tL = BuildDecisionTree(LL)
            tR = BuildDecisionTree(LR )
   end if
   return t
end function
```

Code adapted from Gilles Louppes

stopping criterion
Gini (im)purity = 0

Gini (node L) =

$1 - \sum f(i)^2$

where f(i) is the frequency of the i-th class

Gini (splits Lt and Lr) =

$L_l/L * (1 - \sum f(i)^2) + L_R/L * (1 - \sum f(i)^2)$

where f(i) is the frequency of the i-th class

Note: **splits happen along features!**

Our first example will be a supervised classification problem, in which we are trying to decide if a star is a variable star (RR Lyrae), based on imaging data in 5 bands (u, g, r, i, z) and four colors (u-g, g-r, r-i, i-z)

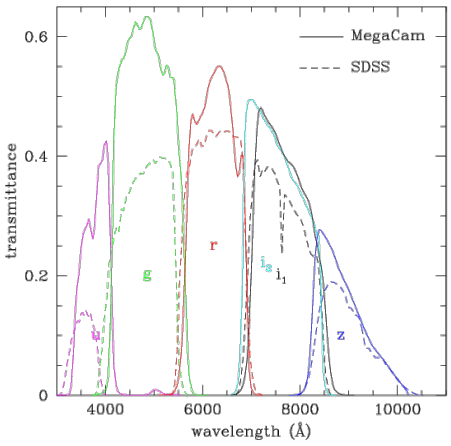

range of observed brightness

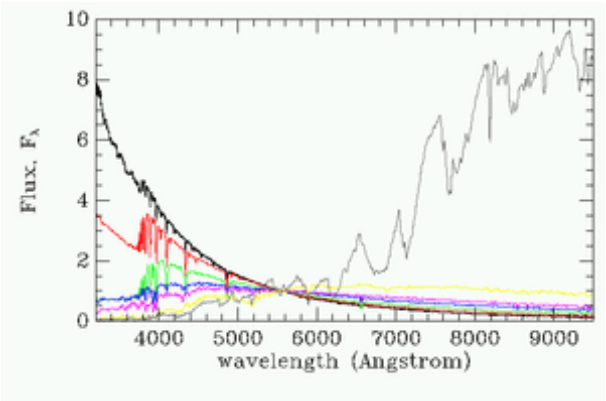

spectra of different stellar types